

5.3 Hello World

A "Hello World" program isn't going to be much different than the one covered in [5.2 FunctionCall](#). There are many ways to make program print "Hello World!" in C++, I will be showing the two main ways. One with `printf()` and another with `std::cout`. It sounds simple, but we will cover some very important concepts.

Source Code

Using `printf()`:

```
1  #include <iostream>
2  int main() {
3      printf("Hello World!\n");
4      return 0;
5  }
```

Using `std::cout`:

```
1  #include <iostream>
2  int main() {
3      std::cout << "Hello World!" << std::endl;
4      return 0;
5  }
```

Reversing the `printf()` Version

Let's view the compiled version of the programs in a disassembler. **We'll start with the `printf()` version.** I'll be using x64dbg for these programs because it does a good job at representing them. To disassemble a program in x64dbg you can go to File > Open then select the program. You can also just drag and drop the program onto the GUI of x64dbg.

When you first throw the binary into x64dbg you may notice that we aren't in the `main()` function. It turns out there are many things added to our program when it's compiled. Depending on how you have x64dbg setup it may start somewhere different. It might start at the "entry point" or near some security cookie stuff.

- **Entry Point** - The entry point is the true beginning of the program. The entry point is ran *before* `main()`. The reason why `main()` isn't called first is because there are other tasks that need to be done before `main()` is called (such as setting up security cookies).
- **Security Cookies** - The security cookie is used to make sure that the return address of a function hasn't been corrupted before following it. This is done to help mitigate buffer overflows. For us, we can ignore anything to do with security cookies.

Finding `main()`

We want to find the `main()` function because that's where the code we want to reverse is at. So how do we get there? The easiest way is by going to the symbol. Symbols are just names that the disassembler was able to find. These names are typically function names. Go the "Symbols" tab and click on the executable.

Base	Module	Party	Path	Address	Type	Ordinal	Symbol
00007FF7F6170000	helloworld_printf.exe	User	C:\Users\ZOF\Desktop\H	00007FF7F6171324	Export	0	OptionalHeader.AddressOfEntryPoint
00007FF937E10000	apphelp.dll	System	C:\Windows\System32\ap	00007FF7F6182000	Import		RtlCaptureContext
00007FF939D00000	kernelbase.dll	System	C:\Windows\System32\Ke	00007FF7F6182008	Import		RtlLookupFunctionEntry
00007FF93C400000	kernel32.dll	System	C:\Windows\System32\ke	00007FF7F6182010	Import		RtlVirtualUnwind
00007FF93CD40000	ntdll.dll	System	C:\Windows\System32\nt	00007FF7F6182018	Import		UnhandledExceptionFilter
				00007FF7F6182020	Import		SetUnhandledExceptionFilter
				00007FF7F6182028	Import		GetCurrentProcess
				00007FF7F6182030	Import		TerminateProcess
				00007FF7F6182038	Import		IsProcessorFeaturePresent
				00007FF7F6182040	Import		QueryPerformanceCounter
				00007FF7F6182048	Import		GetCurrentProcessId
				00007FF7F6182050	Import		GetCurrentThreadId
				00007FF7F6182058	Import		GetSystemTimeAsFileTime
				00007FF7F6182060	Import		InitializeListHead
				00007FF7F6182068	Import		IsDebuggerPresent
				00007FF7F6182070	Import		GetStartupInfoW
				00007FF7F6182078	Import		GetModuleHandleW
				00007FF7F6182080	Import		RtlUnwindEx
				00007FF7F6182088	Import		GetLastError
				00007FF7F6182090	Import		SetLastError
				00007FF7F6182098	Import		EnterCriticalSection
				00007FF7F61820A0	Import		LeaveCriticalSection
				00007FF7F61820A8	Import		DeleteCriticalSection
				00007FF7F61820B0	Import		InitializeCriticalSectionAndSpinCount
				00007FF7F61820B8	Import		TlsAlloc
				00007FF7F61820C0	Import		TlsGetValue
				00007FF7F61820C8	Import		TlsSetValue
				00007FF7F61820D0	Import		TlsFree
				00007FF7F61820D8	Import		FreeLibrary
				00007FF7F61820E0	Import		GetProcAddress
				00007FF7F61820E8	Import		LoadLibraryExW
				00007FF7F61820F0	Import		RaiseException
				00007FF7F61820F8	Import		GetStdHandle
				00007FF7F6182100	Import		WriteFile
				00007FF7F6182108	Import		GetModuleFileNameW
				00007FF7F6182110	Import		ExitProcess
				00007FF7F6182118	Import		GetModuleHandleExW
				00007FF7F6182120	Import		GetCommandLineA
				00007FF7F6182128	Import		GetCommandLineW
				00007FF7F6182130	Import		HeapAlloc
				00007FF7F6182138	Import		HeapFree
				00007FF7F6182140	Import		CompareStringW
				00007FF7F6182148	Import		LCMapStringW
				00007FF7F6182150	Import		GetFileType
				00007FF7F6182158	Import		FindClose

There doesn't appear to be a `main()` function so we will have to find it some other way. One option we have is looking for the string that is going to be printed. We know that our program is going to print "Hello World!" so let's see if we can find where that string is used. We can do this in x64dbg by right-clicking in the disassembly window and going to Search for > All modules > String references. You should see something similar to what is shown in the image below.

Address	Disassembly	String
00007FF70D0E1125	lea rdx,qword ptr ds:[7FF70D0E2260]	"Hello World!\n"
00007FF70D0E1D60	or qword ptr ds:[<__memcpy_nt_iters>],FFFFFFFFFFFFFFFF	"/ "
00007FFB89C31AE8	lea rcx,qword ptr ds:[7FFB89CD12F8]	L"NULL"
00007FFB89C31B9A	lea rax,qword ptr ds:[7FFB89CD12F8]	L"NULL"
00007FFB89C31C1C	lea rax,qword ptr ds:[7FFB89CD12F8]	L"NULL"
00007FFB89C31D4F	lea r8,qword ptr ds:[7FFB89CD12F8]	L"NULL"
00007FFB89C31D86	lea rdx,qword ptr ds:[7FFB89CD12F8]	L"NULL"
00007FFB89C34C52	lea r14,qword ptr ds:[7FFB89CDE500]	L"\$&"
00007FFB89C34D22	lea rbp,qword ptr ds:[7FFB89CDE4B0]	L"\$&"
00007FFB89C3623C	lea rcx,qword ptr ds:[7FFB89CD1410]	"ntdll.dll"
00007FFB89C36280	lea rcx,qword ptr ds:[7FFB89CD1420]	"kernel32.dll"
00007FFB89C36468	lea rcx,qword ptr ds:[7FFB89CD1430]	L"atcuf64.dll"
00007FFB89C36F13	lea r9,qword ptr ds:[7FFB89CE16F8]	L"jj"
00007FFB89C36F47	lea r9,qword ptr ds:[7FFB89CE16E8]	L"&"
00007FFB89C37AAA	lea rcx,qword ptr ds:[7FFB89CD1458]	L"\\??\AtcComm"
00007FFB89C391DE	movzx edx,word ptr ds:[7FFB89CDE960]	L" \"
00007FFB89C3920A	lea rdx,qword ptr ds:[7FFB89CDE960]	L" \"
00007FFB89C3AF86	lea r8,qword ptr ds:[7FFB89CD12E8]	")#"
00007FFB89C3B0E6	lea r8,qword ptr ds:[7FFB89CD12E8]	")#"
00007FFB89C3B694	lea rcx,qword ptr ds:[7FFB89CD1650]	L"ws2_32.dll"
00007FFB89C3B6E8	lea rdx,qword ptr ds:[7FFB89CD1668]	"WSAGetLastError"
00007FFB89C3ECDA	lea rdx,qword ptr ds:[7FFB89CD12F8]	L"NULL"
00007FFB89C3ED11	lea rcx,qword ptr ds:[7FFB89CD12F8]	L"NULL"
00007FFB89C43812	lea rdx,qword ptr ds:[7FFB89CD1D70]	L"UxSubclassInfo"
00007FFB89C43841	lea rdx,qword ptr ds:[7FFB89CD1D90]	L"CC32SubclassInfo"
00007FFB89C438D5	lea rdx,qword ptr ds:[7FFB89CD1DB8]	"UxSubclassInfo"
00007FFB89C438FC	sub rcx,qword ptr ds:[7FFB89CD1DC8]	"CC32SubclassInfo"
00007FFB89C43909	sub rcx,qword ptr ds:[7FFB89CD1DD0]	"lassInfo"
00007FFB89C4B02F	lea rcx,qword ptr ds:[7FFB89CD1DE0]	L"user32.dll"
00007FFB89C4B08D	lea rdx,qword ptr ds:[7FFB89CD25C0]	"IsWindowVisible"
00007FFB89C4B0E6	lea rdx,qword ptr ds:[7FFB89CD25D0]	"GetWindowInfo"
00007FFB89C4B144	lea rdx,qword ptr ds:[7FFB89CD25E0]	"GetWindowThreadProcessId"
00007FFB89C4B196	lea rdx,qword ptr ds:[7FFB89CD26D0]	"WindowFromDC"
00007FFB89C9A9AF	add byte ptr ds:[7FFB89CCC9DA],d1	"MADD2315S"
00007FFB89C9AA2F	add byte ptr ds:[7FFB89CCCA58],d1	"VFMADDSUB213PS"
00007FFB89C9B02F	add byte ptr ds:[7FFB89CCD05A],d1	"X2"
00007FFB89CD0CFA	xor dword ptr ds:[7FFB89CD0D00],esi	"dr0"

As you can see there are quite a few strings. Most of these strings aren't important. Let's search for our "Hello World!" string using the search bar at the bottom. Go ahead and double click on the "Hello World!" string. If there are more than one just choose the first one. This will bring you into the function where the string is used.

00007FF7F6171070 <hellow	\$ 48:83EC 28 SUB RSP, 0x28	sub_7FF7F6171070
00007FF7F6171074	48:8D0D 15 LEA RCX, QWORD PTR DS:[0x7FF7F6189D90]	00007FF7F6189D90:"Hello world!\n"
00007FF7F617107B	E8 90FFFFFF CALL <helloworld_printf.sub_7FF7F6171010>	printf
00007FF7F6171080	33C0 XOR EAX, EAX	
00007FF7F6171082	48:83C4 28 ADD RSP, 0x28	
00007FF7F6171086	C3 RET	

Quick Tip: If you expand the address column you can see the function name. Functions are also marked with a dollar sign ("\$\$") at the start of them.

The function this string is used in is `sub_7FF7F6171070`. This is probably different for you. You can see this next to the address (you might have to resize the column). The name is weird because the disassembler couldn't figure out its name so it just gave it a random name. So how do we know that this is the `main()` function? We don't. That's why reversing can be so difficult. Because we know what the code does and we have access to the source code we know that this is most likely the main function. In a realistic scenario, we wouldn't actually know for sure.

Just so we don't have to keep doing this let's change its label to `main` by right-clicking the first instruction in the function (the one with the `sub_7FF7F6171070` next to it) and choose Label > Label current address. Now we can get back to it by going to the functions tab and searching for `main`. You can get to the functions tab by going to View > Functions or by clicking the "fx" button on the toolbar above the "Memory Map" tab. The toolbar also has the debug buttons (blue arrows).

Analyzing Main

Let's figure out what `main()` is doing by looking at it in Assembly.

00007FF7F6171070 <hellow	\$ 48:83EC 28 SUB RSP, 0x28	sub_7FF7F6171070
00007FF7F6171074	48:8D0D 15 LEA RCX, QWORD PTR DS:[0x7FF7F6189D90]	00007FF7F6189D90:"Hello world!\n"
00007FF7F617107B	E8 90FFFFFF CALL <helloworld_printf.sub_7FF7F6171010>	printf
00007FF7F6171080	33C0 XOR EAX, EAX	
00007FF7F6171082	48:83C4 28 ADD RSP, 0x28	
00007FF7F6171086	C3 RET	

- `SUB RSP, 0x28` - On the first line we can see `RSP` is subtracted by `0x28`. This is part of the function prologue. The function prologue sets up a function and you can usually ignore them.
- `LEA RCX, QWORD PTR DS:[0x7FF7F6189D90]` - `LEA` is short for **Load Effective Address** which will load an address into the destination operand. In this case, the address of our string is being put into `RCX`. Notice that the value is being put into `RCX` which is the first register used to pass parameters to functions. Sure enough, the next line is a function call. What is the `QWORD PTR` part? This is simply stating that the value is a `QWORD` (8 bytes) and it's a pointer. `DS` is short for data segment which can store data that doesn't fit in a register. You can ignore this and just know that it's some data from somewhere in memory. In our case, this address is the address of our string. If you want to view the string in memory you can do so by right-clicking the instruction and going to Follow in Dump > Constant: (Address).
- `CALL <helloworld_printf.sub_7FF7F6171010>` This will call the function at `sub_7FF7F6171010`. As we previously discovered, this function will receive one parameter which is the address of our string. Thanks to xAnalyzer a comment was added that this function is `printf()`. If xAnalyzer didn't figure this out then how would we know this was `printf()`? Typically I would use multiple tools and one of them would probably tell you. The next easiest way would be to debug the program and see what happens

after the call. We would see that the call has something to do with the text being printed out on the screen. Finally, we could reverse the function and figure it out that way.

Now we get to the function epilogue which is the end of the function. This portion will clean everything up before returning execution back to where it was called from.

- `XOR EAX, EAX` - If you XOR something with itself it sets it to zero. In this case, EAX is being zeroed out. I'm not sure why it's zeroing EAX instead of RAX, maybe RAX was never needed and only EAX was. I could dig into this but there really isn't any need to. We can determine from this that the function is returning 0. Based on this we can guess that the function type is an integer. If we look at our source code this is correct. Also, this further solidifies our guess that this is the `main()` function because `main()` is usually an integer type in C/C++.
- `ADD RSP, 0x28` - This is setting the stack to what it was before the function call. Notice that this is opposite to the `SUB RSP, 0x28` at the start.
- `RET` - Returns execution to the caller. This is the end of the function.

I want to touch on the fact that we passed the *address* of the string to `printf()`. In C/C++ strings are just an array of characters. The end of a string is denoted by a null character/byte (0x00). This is known as the string/null terminator. This is important because something like `printf()` will use this to know where the end of the string is so it doesn't print past it. In C/C++ a string is just a character array with the last character being a null terminator.

0x00 is referred to by many names. Null byte, terminator, string terminator, null terminator, and more.

I hope that wasn't too difficult to follow and understand. Finding the `main()` function can be annoying and it certainly doesn't make a good impression for people who are just starting to learn reversing, but it's what we have to do. Feel free to explore a little bit!

Reversing the `std::cout` Version

This is where things get interesting. `std::cout` is part of C++, whereas `printf()` is part of C. This means that the compiler has full control over `std::cout` but not as much control over `printf()`. Think about how you would make a program that just prints text to the screen *once* as efficient as possible. Think about this, what could you do instead of calling a function just once? The answer is to not call any function, just put the text to be printed inside the print function or put the print function with the text to be printed. Sounds really fun for us reverse engineers doesn't it?!

Here is an example:

```
int Add(int num1, int num2){
    return num1+num2;
}
int main(){
    int total = Add(2, 3);
}
```

How could we make this more efficient? There are two ways.

The example below is more efficient because there are no parameters to be passed.

```
int Add(){
    return 2+3;
}
int main(){
    int total = Add();
}
```

The example below is more efficient because there is no function call at all. Instead, the code for the `Add()` function has essentially been copied and pasted into `main()`. It may not seem like it because I have such a small example but that's exactly what has happened in the example. This is called (depending on who you ask) **function inlining, inline expansion, or just in inlining**. I'll just call it inlining.

```
int main(){
    int total = 2 + 3;
}
```

Reversing Time

Open `HelloWorld_Cout.exe` in `x64dbg` and let's get started. First, we want to find the main function. Let's use the same technique we used previously which was searching for the string.

00007FF7467F2134	> 41:8BD0	MOV EDX, R8D	[Arg1
00007FF7467F2137	48:8B01	MOV RAX, QWORD PTR DS:[RCX]	rax:KernelBaseDllInitialize
00007FF7467F213A	FF50 18	CALL QWORD PTR DS:[RAX + 0x18]	sub_[rax+18]
00007FF7467F213D	> 83F8 FF	CMP EAX, 0xFFFFFFFF	
00007FF7467F2140	> 0F84 8A00	JE helloworld_cout.7FF7467F21D0	
00007FF7467F2146	48:FFCF	DEC RDI	
00007FF7467F2149	48:8B0D F0	MOV RCX, QWORD PTR DS:[0x7FF746826240]	
00007FF7467F2150	^ EB A0	JMP helloworld_cout.7FF7467F20F2	
00007FF7467F2152	> 48:6341 04	MOVSDX RAX, DWORD PTR DS:[RCX + 0x4]	rax:KernelBaseDllInitialize
00007FF7467F2156	48:8B4C30	MOV RCX, QWORD PTR DS:[RAX + RSI + 0x48]	[Arg1 = rax+rsi*1+48:KernelBaseDllInitialize+4
00007FF7467F215B	48:8B01	MOV RAX, QWORD PTR DS:[RCX]	rax:KernelBaseDllInitialize
00007FF7467F215E	41:B8 0C00	MOV R8D, 0xC	Arg3 = '\f'
00007FF7467F2164	48:8D15 05	LEA RDX, QWORD PTR DS:[0x7FF746820870]	Arg2 = "Hello World!"
00007FF7467F2168	FF50 48	CALL QWORD PTR DS:[RAX + 0x48]	sub_[rax+48]
00007FF7467F216E	48:83F8 00	CMP RAX, 0xC	rax:KernelBaseDllInitialize, C: '\f'
00007FF7467F2172	> 75 5C	JNE helloworld_cout.7FF7467F21D0	
00007FF7467F2174	> 48:85FF	TEST RDI, RDI	
00007FF7467F2177	> 7E 60	JLE helloworld_cout.7FF7467F21D9	
00007FF7467F2179	48:8B05 C0	MOV RAX, QWORD PTR DS:[0x7FF746826240]	rax:KernelBaseDllInitialize
00007FF7467F2180	48:6348 04	MOVSDX RCX, DWORD PTR DS:[RAX + 0x4]	rax+4:KernelBaseDllInitialize+4
00007FF7467F2184	44:0FB6443	MOVZX R8D, BYTE PTR DS:[RCX + RSI + 0x58]	rcx+rsi*1+58:"m cannot be run in DOS mode.\r\
00007FF7467F218A	48:8B4C31	MOV RCX, QWORD PTR DS:[RCX + RSI + 0x48]	
00007FF7467F218F	48:8B41 40	MOV RAX, QWORD PTR DS:[RCX + 0x40]	rax:KernelBaseDllInitialize
00007FF7467F2193	48:8338 00	CMP QWORD PTR DS:[RAX], 0x0	rax:KernelBaseDllInitialize
00007FF7467F2197	> 74 24	JE helloworld_cout.7FF7467F21BD	
00007FF7467F2199	48:8B51 58	MOV RDX, QWORD PTR DS:[RCX + 0x58]	rcx+58:"am cannot be run in DOS mode.\r\r\n\$"
00007FF7467F219D	8B02	MOV EAX, DWORD PTR DS:[RDX]	
00007FF7467F219F	85C0	TEST EAX, EAX	

What is our string doing there? By looking at the code we can guess that the next function to be called is going to result in the string being printed. Also, note that the first parameter passed to the function is something to do with "KernelBaseDLLInitialize". This is an indicator that it's a function of a library, not a function implemented in `HelloWorld_Cout.exe`. If we are right, this could help us figure out what's going on.

Let's set a breakpoint on the `LEA` instruction that is dealing with our string. You can set a breakpoint by pressing `F2` or by right-clicking and going to Breakpoint > Toggle. Then run the program until you hit that breakpoint. We can then step over the call by clicking the corresponding blue arrow at the top or by pressing `F8`. When we jump over this call we can see that the string appears in the console. It's extremely likely at this point that we are correct that this is `std::cout` or something to do with it. But how do we know that the function that was just called doesn't do something else? You can investigate the function that printed the text into the terminal. The function is accessed via an offset from `RAX`. This offset is `0x48`. Because this function

isn't be called with an address, we will need to know what RAX is and calculate the offset if we want to find the function. So we will have to run the function again until we reach that point. When execution has gotten to that line we can now see what RAX is. Now we go to RAX + 0x48 either by doing our own math or by pressing Enter and letting x64dbg figure it out then follow the call for us. From here you can reverse the function.

Another method to figure out what's going on is to write your own program that uses `std::cout` and see if things match. You could use the graph view to see the layout of the function in the program you're reversing and your own program to see if they match. You could, of course, look at the individual Assembly instructions but I think this is a little more tedious. This is a common technique that more experienced reversers use. Eventually, you will see patterns like this and be able to figure out pretty quickly what's going on.

Getting To Main

So was the string put into `std::cout` or was `std::cout` put into `main()`? Figuring this out is pretty easy. First, go back to the instruction that referenced the string. We want to find out what called this function. Go to the first instruction in this function, this is also the functions address/location. Right-click it and select Find references to > Selected Address(es). We can see that this function is called once, go to where it is called from. This is where it was called from:

00007FF7467F1D80 <helloworld	48:83EC 28 SUB RSP, 0x28	sub_7FF7467F1D80
00007FF7467F1D84	E8 970200 CALL <helloworld.cout.sub_7FF7467F2020>	
00007FF7467F1D89	48:8BC8 MOV RCX, RAX	[Arg1 = rax:KernelBaseDllInitialize
00007FF7467F1D8C	E8 3F0500 CALL <helloworld.cout.sub_7FF7467F22D0>	sub_7FF7467F22D0
00007FF7467F1D91	33C0 XOR EAX, EAX	
00007FF7467F1D93	48:83C4 28 ADD RSP, 0x28	
00007FF7467F1D97	C3 RET	

This appears to be `main()` because of how small it is, but there are no guarantees. This is the function that causes our string to be printed so it's very likely that this is `main()` or something related to it.

So we now know that in this case, the compiler decided to put our string into `std::cout`.

What's with the second call? I encourage you to find out what this call is on your own. Here is a quick hint: the two functions called are very similar.

► Answer (click to reveal):

The second call is `std::endl`. You will see that the function prints "\n" which is the newline character.

`std::cout` Used Twice

As I said, inlining is almost never used when a function is called more than once. So what does it look like if `std::cout` is used twice? Well, it look much more like the `printf()` version.

00007FF749531D80 <testing	48:83EC 28 SUB RSP, 0x28	sub_7FF749531D80
00007FF749531D84	48:8D15 E8 LEA RDX, QWORD PTR DS:[0x7FF749560870]	00007FF749560870: "Hello World!"
00007FF749531D8B	E8 B00200 CALL <testing.sub_7FF749532040>	
00007FF749531D90	48:8BC8 MOV RCX, RAX	[Arg1 = rax:KernelBaseDllInitialize
00007FF749531D93	E8 680500 CALL <testing.sub_7FF749532300>	sub_7FF749532300
00007FF749531D98	48:8D15 E8 LEA RDX, QWORD PTR DS:[0x7FF749560880]	[Arg1 = "Hello Again!"
00007FF749531D9F	E8 9C0200 CALL <testing.sub_7FF749532040>	sub_7FF749532040
00007FF749531DA4	48:8BC8 MOV RCX, RAX	[Arg1 = rax:KernelBaseDllInitialize
00007FF749531DA7	E8 540500 CALL <testing.sub_7FF749532300>	sub_7FF749532300
00007FF749531DAC	33C0 XOR EAX, EAX	
00007FF749531DAE	48:83C4 28 ADD RSP, 0x28	
00007FF749531DB2	C3 RET	

As you can see there is one call for `std::cout` and one call for `std::endl` for each one of the two strings being printed, totaling in four function calls.

More About Finding `main()`

There are a few techniques that you can use to find `main()`. In the future, I will talk about these methods. One thing I did want to mention is that `main()` should be called in the entry function. If you're looking for `main`, it's best to start from the bottom of the entry function and work your way up. That's it for now, I will add more about this topic in the future.

Final Notes

The `printf()` version should be fairly easy to understand. The `std::cout` is much more confusing. Hopefully, you get the general idea of what's going on though. From this point, I would highly recommend that you write your own programs and see what's going on. Visual Studio allows you to view the Assembled version of your program along with the source code which can be great for learning. Don't worry if this bit was confusing, the `std::cout` thing is really weird and counter-intuitive.

[<- Previous Lesson](#)

[Next Lesson ->](#)

[Chapter Home](#)